

Machine Translation Evaluation

Machine Translation Evaluation

- Human subjective evaluation is the best but is time-consuming, expensive and inconsistent for long sentences.
- Automated evaluation comparing the output to **multiple human reference translations** is cheaper and generally correlates with human judgements.

The BLEU Score

- BLEU has been shown to correlate with human evaluation when comparing outputs from different SMT systems
- However, it does not correlate with human judgments when comparing SMT systems with human translations
- Other MT evaluation metrics have been proposed that claim to overcome some of the limitations of BLEU

Other Evaluation Methods

- Direct Assessment (Graham et al., 2015): a protocol for evaluation by crowd-sourcing
- We will discuss more nuanced semantic evaluation later in the course

Yvette Graham, Timothy Baldwin, Nitika Mathur.

"Accurate Evaluation of Segment-level Machine Translation Metrics", NAACL 2015